
PRODUCT RESEARCH · 27 AUGUST 2026

Eight businesses built on measured truth

Eight industries where the numbers that decide money are quietly guessed — and a product shape that measures them, marks what it does not know, and carries the result all the way to a priced, signed document.

Sourced competitor pricing · ranked failure modes · one settling test per idea

8 briefs

Contents

01	The eight at a glance
02	The test every idea had to pass
03	1 — Templating
04	2 — Fire & life-safety inspection
05	3 — The binding estimate
06	4 — The honest weight
07	5 — Range of motion
08	6 — Floor prep
09	7 — Barrier survey
10	8 — What's behind this wall
11	Sources

How to read the sourcing

Every figure in this document carries where it came from. Nothing is asserted without one of these four marks.

Mark	Means
VENDOR	Read off the company's own page.
REVIEW	Reported by a third-party review or comparison site and not independently confirmed . Several such sites are run by a competitor in the category they price. Verify before acting.
STANDARD	A published technical or legal standard.
ASSUMPTION	Mine, marked at every point of use.

No figure in this document was produced on a device I own, and no accuracy claim in it was measured. Where an unmeasured number decides whether an idea works, the brief says so and says it is unmeasured. The closing pages list every figure deliberately not claimed.

The eight at a glance

Eight industries where a number that decides money is currently guessed.

Each one is a product that **measures** that number, **marks what it does not know**, lets a human override the sensor and resolves around the override, and carries the result all the way to a priced, signed document — sold flat and unlimited into a field that meters, with no server behind it.

The eight

#	Idea	Industry	The wedge, in one line
1	Templating	Stone / countertop fabrication	Refuses to export to CNC until a tape has touched every cut edge
2	Fire & life-safety inspection	Fire protection	Works in the basement; the report says what was <i>assumed</i> , not just what passed
3	The binding estimate	Household goods moving	The estimate prints what was measured and what was guessed
4	The honest weight	Cattle production	The only one that publishes its error band instead of claiming $\pm 3\%$
5	Range of motion	Physical therapy / chiropractic	No server means patient data never reaches infrastructure you own
6	Floor prep	Flooring installation	Won't certify the slab — tells you where to put the straightedge, then prices the prep
7	Barrier survey	Accessibility compliance	A survey is evidence, and evidence has provenance or it is worthless
8	What's behind this wall	Construction / property	Sells the homeowner a permanent record — a new invoice line, not a cost

Scored on six axes

Innovation, industry need, competitive advantage, buildability, creativeness, monetisation. **1–5 each. These scores are judgement, not measurement** — they exist to force a ranking, and the reasoning behind every one is in the brief that follows.

#	Idea	Innov.	Need	Advantage	Build	Creative	Money	Total
7	Barrier survey	4	4	5	4	4	5	26
8	What's behind this wall	5	3	5	4	5	4	26

#	Idea	Innov.	Need	Advantage	Build	Creative	Money	Total
1	Templating	4	5	4	5	3	4	25
6	Floor prep	3	5	4	5	3	3	23
2	Fire & life-safety	3	5	4	2	3	5	22
3	The binding estimate	4	4	3	3	4	4	22
5	Range of motion	4	4	4	3	4	3	22
4	The honest weight	3	3	2	3	5	2	18

Buildability is scored against a hard constraint — one person, iOS plus web, offline-first, no server, zero marginal cost per customer. That constraint cut more candidates than the market analysis did. It is why there is no marketplace here, no two-sided network, and nothing that depends on a licensed third-party database.

Which to build

7 — the barrier survey — is the strongest business. The buyer already bills **\$1,700–\$4,500 per survey**, so the return on a subscription is a single line of arithmetic rather than an argument. And a barrier survey is read in litigation, which makes "what was measured versus what was assumed" a professional necessity rather than a nice-to-have. Its risk is distribution, not product: outside California there is no register of accessibility consultants to sell into.

8 — what's behind this wall — is the best idea. It is the only one of the eight where the contractor **makes** money instead of spending it: the record is a \$150–\$400 line item on a job where the walls were already open, which is a far shorter sale than any subscription pitch. It is also the most defensible — nobody in the category has noticed that the person who benefits most is the homeowner, not the contractor. Its risk is that the value arrives four years late, and latent value has to be sold rather than bought.

1 and 6 are the fastest to a working prototype, and both are gated on a single unmeasured accuracy figure that one afternoon of testing would settle.

4 is the one not to build. It is the most creative idea here and it is genuinely late: four phone-based competitors already exist, one of them already using the same sensor.

Before writing any code

Every brief ends with **one test that settles it**, each chosen to cost a day or less and to be runnable *before the first commit*. Idea 1's test costs one slab. Idea 6's costs one afternoon. Idea 8's costs five phone calls.

Several of these ideas are gated on an accuracy figure that **does not exist in public sources and was not measured here**. Those gates are named in each brief and listed together in the closing pages. Running the test is cheap. Building first and discovering the sensor cannot hold the tolerance is not.

The test every idea had to pass

Eight ideas. Each one is the answer to a single question: *what else looks like Trueline?*

Not "what else is a good app." Trueline is a specific shape, and the shape is what was copied.

1. What Trueline actually is, read out of its own repository

Five properties, all of them read out of `README.md` , `DECISIONS.md` , `docs/BUSINESS.md` and `docs/money.md` rather than from the name:

- 1. **A sensor measures, a human overrides, and the model re-solves around the override.** Type one wall's real width and every connected room, corner and quantity moves with it. Competitors let you edit a number without re-solving.
- 2. **Every figure carries its provenance** — `scanned` , `measured` , `derived` , `adjusted` — on screen and on every export. The app refuses to issue a drawing from numbers a tape has never touched.
- 3. **The whole job, not a slice.** The takeoff, the proposal, the change order and the invoice all come off the same measured walls. Two products next to each other cannot do that at any price.
- 4. **Flat and unlimited, against a field that meters.** Per scan, per project, per space — metering punishes the low-volume operator, and almost every remodeler is a low-volume operator.
- 5. **No server.** Storage is the customer's own iCloud, marginal cost per customer is zero, and the thing works in a basement with no signal.

2. The four patterns, as a gate

Every idea below had to clear **all four**, not three:

Pattern		The question asked of each idea
A	Trusted-data engine	Does this industry quietly guess, and does everyone hide it?
B	Whole job vs. pieces	Do 4–6 incumbents each own a slice, with the buyer retyping between them?
C	Sensor → money	Does phone hardware capture the real world and carry it to a priced, signed document?
D	Underserved small operator	Does everyone meter, and does that punish the one-truck shop?

And then the build constraint, which killed more candidates than the patterns did:

Solo-buildable. iOS + web. Offline-first. No server. Zero marginal cost. App Store subscription.

That constraint is why there is no marketplace here, no two-sided network, no "AI copilot for X" that dies the moment an API key expires, and nothing that needs a licensed parts database. Auto-body collision estimating was cut for exactly that

reason: CCC, Mitchell and Audatex license the parts and labour data, and a licensed database is a per-customer cost, which breaks the model.

3. What is sourced and what is not

The same discipline `MARKET-RESEARCH.md` uses, because a brief without it is a pitch.

- **VENDOR** — read off the company's own page.
- **REVIEW** — reported by a third-party review, comparison or industry site and **not independently confirmed**. Comparison sites are often run by a competitor or an affiliate; treat every one of these as a lead to verify, not a fact.
- **STANDARD** — a published technical or legal standard.
- **ASSUMPTION** — mine. Marked every time.

Researched 2026-08-27. **No figure below was produced on a device I own, and no accuracy claim in this document was measured.** Where a number would decide whether an idea works, it says so and says it is unmeasured.

Every URL is in `../SOURCES.md`.

4. The honesty that costs something

Two of the eight are weaker than they first looked, and both say so in their own brief rather than in a footnote:

- **Idea 4 (livestock) is contested.** At least four phone-based cattle weight products already exist. It stays in the set because of *how* they compete, not because the space is empty — but it is the riskiest of the eight.
- **Idea 6 was wrong as first conceived.** The plan was a phone that produces an ASTM E1155 floor-flatness F-number. It cannot: the standard needs precision a phone does not have. The idea survives only in a reduced, honest form, and the reduction is written up rather than quietly dropped.

That is the same move `DECISIONS.md` makes when it records that the DXF recommendation was wrong and corrects it in place. A research file that never says "this was wrong" is not being used.

1 — Trueline for the fabrication shop

Digital templating for countertops, stair treads and flooring, on a phone instead of a \$25,000 rig.

A template is the shape of a real, crooked opening, captured accurately enough to cut stone against it. Get it wrong by 3 mm and you have scrapped a \$2,000 slab.

The job today

A small fabrication shop has three options, and all three are bad.

- 1. **Cardboard and hot glue.** Free, slow, and it warps in a hot truck.
- 2. **A digital templating rig.** Accurate, and priced like a car.
- 3. **Subcontract the templating** to somebody who owns a rig, and pay per job forever.

The rigs: Prodim Proliner is reported at **\$20,000–\$45,000** [REVIEW](#), with one review site putting the common configuration at **\$22,000–\$30,000** [REVIEW](#). Laser Products' LT-2D3D — described as the most widely used system in US countertop shops — is reported at **\$18,000–\$25,000** [REVIEW](#). The older LT-55 turns up used at **\$3,599–\$8,995** [REVIEW](#), which tells you what the second-hand market thinks a decade-old rig is worth.

That price is the whole market structure. A shop doing four kitchens a month cannot justify \$25,000 of capital, so it stays on cardboard, or it subcontracts, or it does not bid the job.

Where it sits

Product	What it is	Price	Source class
Prodim Proliner	Digital templating hardware, stone industry	\$20,000–\$45,000	REVIEW
Laser Products LT-2D3D	Laser templating, largest US installed base	\$18,000–\$25,000	REVIEW
Laser Products ScanTemplater	Camera/photogrammetry alternative	not published	REVIEW
Flexijet 3D, Leica 3D Disto	Complex 3D work	not published	REVIEW
Polycam, Scaniverse, Scandy Pro	General phone LiDAR, 3D mesh out	~\$150/yr Basic (Polycam)	REVIEW
This	Phone template → verified outline → CNC file → priced quote	flat subscription	—

The gap, stated precisely. General phone-LiDAR apps produce a *mesh*. A fabricator does not need a mesh; they need a **closed 2D outline with exact lengths and angles, and a seam plan**. Nobody sells the phone-priced version of the

templating job — the search for one returns 3D scanner apps that "aren't specifically designed for countertop templating workflows" REVIEW.

The wedge

A phone is not accurate enough to template a countertop, and that is the product.

This is Trueline's argument, transplanted whole. The scan gets the outline and the rough geometry. Then the app hands back **the four measurements worth a tape** — ranked, exactly like Trueline's carry list — and the fabricator types the real numbers. The outline re-solves around them: type the true back-wall length and the sink cut-out, the overhang and the seam all move with it.

Every dimension exports marked `scanned` or `measured`. **A template with any scanned dimension on a cut edge will not export to CNC.** The app refuses.

That refusal is the sales demo. Nobody scraps a slab because a tape was wrong. They scrap it because a number nobody checked looked exactly like a number somebody had.

The four-pattern check

- **A — trusted data.** A cut edge is pass/fail at ± 1.5 mm. Every general scanning app hands over a number with no provenance.
- **B — whole job.** Templating rig → drawing software → quoting → CNC post is four things today. Here the square footage that priced the job is the same polygon that drives the cutter.
- **C — sensor → money.** Outline → sq ft → priced quote with edge profile and cut-out charges → signed → DXF to the saw.
- **D — small operator.** The gate is a \$25,000 capital purchase, which is the most brutal metering there is: not per-scan, per-shop.

What gets built

Most of it exists in Trueline already:

Needed	State
Exact integer geometry, nanometre lengths	Built — <code>core/</code>
Feet-inches-fractions parsing and formatting	Built
Provenance model (<code>scanned</code> / <code>measured</code> / <code>derived</code>)	Built
Re-solve around a typed measurement	Built — the rectilinear solver
Angled walls held fixed while square ones absorb error	Built — exactly what a chamfered counter needs
DXF with dimensions	Built

Needed	State
Ranked "worth a tape" carry list	Built
Takeoff and price book	Built
Seam planning, edge profiles, cut-out library	New
Slab yield / nesting	New, and the hardest part

Realistically 30–40% of this is already compiled and passing tests in Trueline. That is the strongest buildability score in the set.

The gate and the price

- **Free forever** — template, outline, dimensions, 2 templates kept at once.
- **Paid** — CNC export, seam plan, priced quote, unlimited templates.

\$59/month, flat, unlimited. Above a general scanning app, far below the \$18,000 floor of the alternative. The comparison a fabricator makes is not against software; it is against a capital purchase they have already decided they cannot make.

Where it fails, ranked by likelihood

1. **Accuracy, and it is not close to a footnote.** (*Most likely.*) A fabricator's tolerance is roughly ± 1.5 mm on a cut edge. **No accuracy figure for phone LiDAR templating exists in this research, and none was measured.** If the human-verified path cannot hold that tolerance in a real kitchen, this is not a product at any price. *The tell:* the first scrapped slab.
2. **Templating is a trade, not a task.** (*Plausible.*) Rig operators are trained. A phone that makes templating look easy may produce confident amateurs, whose scrapped slabs become your App Store reviews.
3. **The shop with a rig will not switch,** and the shop without one may not be a buyer at all.

The one test that settles it

Template one real kitchen with a phone and the human-verified path. Cut it. Does it fit? That test costs one slab and about a day, and it decides the whole idea before a line of the new code is written.

Do that test before building anything.

2 — The inspection that has to hold up

Fire and life-safety inspection, done offline in a mechanical room, ending in a signed report and a priced repair quote.

An NFPA inspection report is not a work product. It is a **legal document** that an authority having jurisdiction, an insurer and eventually a lawyer will read.

The job today

A one-truck fire-protection contractor inspects sprinkler systems, extinguishers, alarms and backflow preventers. The work happens in basements, mechanical rooms, stairwells and riser closets — **the places with no signal**. The inspector writes on paper, drives back, and retypes everything that evening.

Then the deficiencies found get retyped a second time into whatever produces the repair quote, and a third time into whatever invoices it.

Where it sits

Product	What it is	Price	Source class
Inspect Point	Fire inspection & service software	not published	VENDOR
BuildingReports / ScanSeries	Barcode-scanned device inspection, compliance reporting	"Contact for pricing"	VENDOR
iWorQ	Fire inspection management, AHJ-side	not published	VENDOR
SafetyCulture	General inspection checklists	published tiers	VENDOR
This	Offline inspection → deficiency → priced repair → signed report	flat, published	—

The finding worth more than any price I could have found: there is no price. Both category leaders quote. Neither publishes a number. I searched for per-device pricing and the result was that both "use custom/quote-based pricing models rather than transparent per-device pricing structures" [REVIEW](#).

A market where nobody publishes a price is a market where the small buyer is being charged what they will bear, and where a **published, flat, unlimited number is itself the marketing**. That is Trueline's flat-vs-metered argument, handed to you by the competition's own pricing page.

The wedge

Three things, in the order a buyer cares:

1. It works in the basement. Offline-first is a checkbox in every competitor's feature list and the actual daily reality of this trade. Trueline is already built full-offline with real conflict resolution on sync. Here that is not a nice-to-have — it is the job.

2. The report says what was tested and what was assumed. A device the inspector could not reach, a valve that could not be exercised, a head hidden above a hard ceiling — today these silently become a clean line on a report. Trueline's provenance model marks each one, and the report prints the distinction. **An inspector's signature is on that document.** Being able to prove what you did and did not verify is not a feature; it is professional indemnity.

3. The deficiency becomes money the same day. Found → photographed with camera pose → priced off the price book → quoted → signed, before leaving the building. Today that is three products and two rounds of retyping, and a meaningful share of deficiencies never get quoted at all because the retyping never happens.

The four-pattern check

- **A — trusted data.** Provenance on a legal document, where "not verified" is a real and currently invisible category.
- **B — whole job.** Inspection software + quoting + invoicing is two or three subscriptions that share no data.
- **C — sensor → money.** Camera with pose, measured travel distances and head spacing, barcode/QR device identity → deficiency → priced repair → e-signature.
- **D — small operator.** Quote-based pricing is metering with the meter hidden.

What gets built

Reuses the Trueline halves almost wholesale: photos-with-pose, the provenance model, the price book, the proposal with e-signature and audit trail, the scheduling and invoicing, the offline sync, the PDF pipeline.

New: the NFPA form logic (25 for sprinkler, 10 for extinguishers, 72 for alarms), device inventory with barcode identity, and inspection *frequency* tracking — annual, quarterly, five-year — which is the recurring revenue engine of the whole trade and the reason a customer never churns.

The frequency tracker is the retention mechanism. A contractor whose next twelve months of inspection obligations live in your app does not cancel in month four. Compare that to Trueline's own stated risk, where churn "is the one number that decides whether this is a business."

The gate and the price

- **Free forever** — inspect, photograph, one building's device inventory.
- **Paid** — reports, deficiency quoting, unlimited buildings, frequency scheduling, e-signature.

\$89/month, flat, unlimited buildings and devices, price published on the page. Being the only published price in the category is the pitch.

Where it fails, ranked by likelihood

1. **The forms are the product, and the forms are enormous.** (*Most likely.*) NFPA 25 alone is a large, versioned, jurisdictionally-varied body of requirements, and AHJs accept different formats. This is a research and maintenance burden that never ends, and it is not a burden a solo builder carries lightly. **This is the single biggest risk in the set.**
2. **Certification and liability.** (*Plausible.*) If the app formats a report an AHJ rejects, the contractor is exposed. Some jurisdictions mandate specific systems outright.
3. **BuildingReports owns the barcode.** (*Plausible.*) Their model puts a scannable tag on every device in a building. A building already tagged is a building already switched — that is a real moat and it is already installed.

The one test that settles it

Take one real NFPA 25 annual sprinkler report that an AHJ has already accepted, and reproduce it exactly from the app's data model. If the form logic for one report type is a week, this is buildable. **If it is a month, it is not** — and you will know inside that first week.

3 — The binding estimate

LiDAR inventory of a house → cubic feet → a binding moving estimate the mover can actually stand behind.

A moving estimate is a promise about a number nobody measured.

The job today

A mover walks a house, eyeballs the contents, and writes a cube sheet from memory and habit. That guess becomes a price. If the guess is low the mover eats the difference or has a fight on the driveway on moving day; if it is high they lose the job to whoever guessed lower.

The entire industry runs on an unmeasured number that is contractually binding. It is difficult to invent a cleaner fit for Trueline's thesis.

Where it sits

Product	What it is	Price	Source class
Yembo	AI video survey → visual inventory → estimate	not published; per-survey model	VENDOR / REVIEW
HomeSurvey.ai	Virtual survey, positioned against Yembo	\$20 per survey, no subscription	REVIEW
SmartMoving	Moving company CRM/ops	from \$399/month flat	VENDOR / REVIEW
MoveitPro	Moving + storage ops	\$800–\$3,000/month	REVIEW
Movegistics	Moving ops	\$1,000–\$4,000/month	REVIEW
This	Measured inventory → cube → binding estimate → bill of lading	flat, ~\$99/mo	—

Two structural facts, both exploitable.

- The survey layer meters, the ops layer does not.** Yembo is per-survey; HomeSurvey.ai is explicitly **\$20 per survey**

REVIEW

. A mover doing 60 estimates a month to win 20 jobs pays for all 60. **The metering is on the losing bids.** That is the most punishing meter placement possible, and it falls hardest on the small mover with the lowest close rate.
- The ops layer starts at \$399/month and runs to \$3,000**

REVIEW

 — and it does not measure anything. So the small mover pays twice, and the two halves share no number.

The wedge

Yembo estimates from video. This one measures, and says which is which.

A video survey infers volume from pixels. LiDAR measures it. The difference is that a measured sofa has a real dimension and an inferred sofa has a confident guess — and on a *binding* estimate, the difference between those two is the mover's margin.

So: scan the room, get the outline and the large items measured. Then the app does the Trueline move — it hands back **the items worth a tape measure**, ranked by how much cube they contribute and how uncertain they are, and the mover types the real numbers. Everything re-solves.

Then the part nobody does: **the estimate prints what was measured and what was estimated**. A customer disputing a moving bill and a mover defending it are looking at the same document, and it tells the truth about its own inputs. The packed-box allowance is an allowance and says so; the piano is measured and says so.

The four-pattern check

- **A — trusted data.** A binding contract built on an eyeballed number, with no provenance anywhere in the category.
- **B — whole job.** Survey tool + CRM + estimate + bill of lading, all separate, all retyped.
- **C — sensor → money.** LiDAR → cubic feet → weight → binding estimate → signed bill of lading. This is the most direct sensor-to-money line of the eight.
- **D — small operator.** Metering on *losing* bids, plus a \$399 floor for the ops half.

What gets built

Trueline's room capture, volume arithmetic, provenance model, price book, proposal-with-signature and offline sync carry over.

New: an item recognition and cube library (sofa, dresser, king mattress → standard cube), the cube-to-weight conversion, packing materials, and the bill of lading. **Item recognition is the risk** — but note that it degrades safely: the app can ask the mover to tap and name items rather than recognise them, which is slower and still faster than a clipboard. It never needs a model that works perfectly, because a human is standing there.

The gate and the price

- **Free forever** — scan, inventory, cube total, 2 surveys kept.
- **Paid** — binding estimate, bill of lading, unlimited surveys, crew sheet.

\$99/month, flat, unlimited surveys. Against \$20/survey, this pays for itself at five estimates a month, and every estimate after that is free. **Lead with that arithmetic** — it is the entire sale, and it is aimed exactly at the small mover the per-survey model punishes.

Where it fails, ranked by likelihood

1. **Movers do not want an accurate number.** (*Most likely, and the deepest risk in the set.*) A binding estimate that is honestly high loses the job to a competitor who guessed low. The industry's tolerance for vagueness may be a *feature* of how movers win work, not a defect they want fixed. **If that is true, no product fixes it.** *The tell:* demos land well with operations, and sales says no.
2. **Yembo is funded and entrenched** with major van lines. Atlas Van Lines already runs a Yembo-powered estimator VENDOR.
3. **The survey is often done by the customer, not the mover** — self-survey is one of Yembo's four modes VENDOR. A customer will not measure anything, and an app that demands verification cannot be handed to them.

The one test that settles it

Ask three small movers one question: "*When your estimate is wrong, who pays?*" If they say the customer, there is no product. If they say they eat it — and they can tell you roughly how much a year — you have found the budget.

4 — The honest weight

Depth scan of an animal → weight estimate → market-priced lot → consignment sheet.

Read the risk section first. This is the most contested idea of the eight, and it is in the set for a specific reason that is not "the space is empty."

The job today

A cattle producer sells by weight. Knowing the weight before the animal is on somebody else's scale decides when to sell, what to feed, and what a lot is worth. The instrument for this is a chute scale, reported at **\$3,000–\$15,000** [REVIEW](#), and running animals through a chute has a real biological cost — stressed cattle "can lose up to 5% of body weight during a single weighing event" [REVIEW](#).

So the small producer does not weigh. They eyeball it, and eyeballing weight is a skill with a wide and unmeasured error band.

Where it sits — and this is the problem

Product	What it is	Claim / price	Source class
CattleWeight AI	Phone LiDAR weight estimation	±3% accuracy , 15-second scan; launching end of 2025	VENDOR
Luzardo	Smartphone photo → AI weight	contactless, seconds	VENDOR
CattleWorth (MJE)	Photo → weight and market value	97% accuracy	VENDOR
Cattle Weight Pro	Tape-measure-assisted estimate	\$1.99 one-time	VENDOR
Marubeni	Trial-launched cattle weight app	—	VENDOR
Chute scales (Tru-Test, Avery, AgriEID)	Physical weighing	\$3,000–\$15,000	REVIEW

Four phone-based products already exist, one of them already does LiDAR, and one already does weight-plus-market-value — which was going to be the wedge.

That is a hard finding and it is stated here rather than buried, because a brief that omitted it would have been worth nothing.

Why it is still in the set

Look at what those vendors are claiming: $\pm 3\%$. 97%.

Now read Trueline's own rule, from `docs/BUSINESS.md` :

Never lead with accuracy. There is no measured accuracy figure. The research says so in writing: no accuracy claim in it was produced on a device we own. Claiming one is the fastest way to be caught.

Every competitor in this market is leading with an accuracy claim, and I could find no independent validation of any of them. Not one published error band by breed, by body condition, by camera angle, by lighting. " $\pm 3\%$ " is a number in marketing copy.

That is a market of confident claims, and confident claims break on contact with a real sale barn — where the buyer's scale is the only number that counts and your app was 40 lb off on a heavy-muscled bull.

The wedge is therefore the inverse of the category's: the product that publishes its own error band, states it per breed and per condition, widens the band when it cannot see the animal properly, and **refuses to produce a consignment figure when the band is too wide to sell on**. Plus the whole job — weight → current market price → lot value → consignment sheet → sale record reconciled against the actual scale ticket, which is the thing that makes the error band *real* instead of claimed. Every sale teaches the app what it got wrong.

That reconciliation loop is the moat. It is also the only honest way anyone will ever produce a validated accuracy figure in this category.

The four-pattern check

- **A — trusted data.** Cleanly met, and by inversion: the whole category is making unvalidated claims.
- **B — whole job.** Weight, market price, lot composition and consignment are separate today. CattleWorth reaches furthest and still stops before the sale record.
- **C — sensor → money.** Depth scan → weight → cwt price → lot value → consignment sheet.
- **D — small operator.** A \$3,000–\$15,000 chute scale is a capital gate on a 30-head producer.

What gets built

Volume-from-depth, the provenance model, offline-first (cattle pens have no signal) and the document pipeline carry over from Trueline. The reconciliation ledger — predicted vs. actual scale ticket — is new and is the core asset.

The hard part is not the app. It is the validation dataset, and it cannot be bought or reasoned into existence. It has to come from producers who scan, sell, and report the real ticket back. That is a slow cold start with no server to coordinate it.

The gate and the price

- **Free forever** — scan, weight estimate with its honest band, herd list.

- **Paid** — market pricing, lot builder, consignment sheets, reconciliation history, unlimited animals.

\$29/month. Agricultural software prices below trade software, and a \$1.99 one-time competitor already anchors the low end.

Where it fails, ranked by likelihood

1. **Four competitors, one funded conglomerate, and you are late.** (*Most likely.*) CattleWeight AI already ships the exact capture technology.
2. **The honesty pitch may not sell here.** (*Plausible, and this is the one that worries me.*) A producer comparing two apps sees " $\pm 3\%$ " next to your " $\pm 8\%$ on this animal, and I will not give you a number on that one." **The honest product looks worse on the shelf.** Trueline has the same problem, but sells to a contractor whose signature is on a document — a producer has no equivalent forcing function.
3. **Cold start.** The reconciliation loop needs volume to be worth anything, and it is worth nothing on day one.

The recommendation

Do not build this fourth. It is here because it clears all four patterns honestly and it is genuinely creative — but it is the only one of the eight where a well-funded competitor already ships the same sensor into the same hand.

If it gets built, build it as an **honest-band engine plus reconciliation ledger**, never as "another cattle weight app," and expect the validation dataset to take a year.

5 — The measurement the claim depends on

Range-of-motion and functional assessment for the solo physical therapist, chiropractor and hand therapist — measured on device, documented for reimbursement, and no patient data ever leaves the phone.

The job today

A physical therapist measures a shoulder with a plastic goniometer, reads it to the nearest five degrees, and writes it into a note. That number justifies the plan of care, supports medical necessity, and is what an insurer reads when deciding whether to pay.

It is the single most consequential number in the encounter, and it is taken with an instrument described in the literature as "over 40 years old" with known "measurement discrepancies" [REVIEW](#).

Then it gets retyped into the EMR.

Where it sits

Product	What it is	Price	Source class
WebPT	PT EMR	from \$99/provider/month; SMS, telehealth and prior-auth billed separately	REVIEW
Jane App	Multi-discipline EHR/PM	from \$54/month; \$100–150/mo typical for a US PT practice with billing	REVIEW
PT EMR category generally	EMR + billing + outcomes	\$150–\$300/provider/month	REVIEW
Halo	Digital goniometer (hardware)	claims accuracy within 1°	VENDOR
GeniusPT, PivotalPT, PhysioMaster	ROM capture apps	various	VENDOR
This	Measured ROM with provenance → billable documentation → claim-ready note	flat, ~\$49/mo	—

Two facts that make this a Trueline shape.

1. **The measurement layer and the money layer are different products.** The ROM app measures; the EMR bills; the therapist retypes between them. Same seam, same retyping, same eleven-at-night as the remodeler with magicplan and Jobber.
2. **The category has a validation problem it admits in print.** "Goniometer apps need to be validated as reliable clinical tools before being integrated into regular physical therapy practice," and of the apps in a 2014 systematic

review, **only five are still on the market** [REVIEW](#). A field littered with abandoned measurement apps is a field where nobody trusted the numbers.

The wedge

Three things, and the third is the one nobody can copy.

1. The vision estimate never overrides the clinician. On-device pose estimation gives a fast ROM figure. The therapist's goniometer reading overrides it and the record re-solves — same rule as a typed wall width. Every value in the note is marked `vision-estimated` or `clinician-measured`. **An insurer auditing the chart can see which is which.** No competitor does this, and it is precisely the thing an auditor cares about.

2. The measurement lands in the note already billable. ROM deficit → documented functional limitation → the note that supports medical necessity. One artefact, not two products.

3. No server, and here that is a regulatory superpower rather than a constraint. Trueline's zero-marginal-cost architecture — everything on device, storage in the customer's own iCloud, AI provider off by default — means **protected health information never transits infrastructure you own.** You are not a business associate for data you never receive. Every competitor is a cloud EMR carrying full HIPAA business-associate exposure, breach liability and the audit burden that follows.

That is the same architectural decision Trueline already made, worth an order of magnitude more in this industry than in construction.

The four-pattern check

- **A — trusted data.** A number an insurer relies on, taken with a 40-year-old instrument, recorded with no provenance.
- **B — whole job.** Measurement app + EMR + billing, retyped between.
- **C — sensor → money.** Camera pose → measured ROM → documented limitation → claim.
- **D — small operator.** \$99–\$300 per provider per month, with the category's own reviewers saying that for a 1–5 provider clinic "the complexity and cost often outweigh the benefits" [REVIEW](#).

What gets built

The provenance model, the override-and-re-solve pattern, offline-first, the PDF pipeline and the document generator carry over. Apple's Vision framework does on-device body pose; no server, no API, no per-patient cost.

Deliberately not built: billing submission. Clearinghouse integration is a server, a per-claim cost and a compliance surface. This produces the **documentation**; it does not become an EMR. That boundary is the same one Trueline draws with Plumblin — a clean seam rather than a swallowed scope.

The gate and the price

- **Free forever** — measure ROM, track a patient's progress, 2 patients kept.
- **Paid** — note generation, export, unlimited patients, re-assessment comparison.

\$49/month per clinician, flat. Half of WebPT's starting price, and it is not trying to be WebPT.

Where it fails, ranked by likelihood

1. **Clinical validation is the entire product and it is expensive.** (*Most likely.*) Selling a measurement into a clinical record means somebody will ask for the validation study. **No accuracy figure for on-device pose estimation against goniometry exists in this research and none was measured.** The literature's own verdict is that these apps need validating first.
2. **Regulatory classification.** (*Plausible, and it is a hard stop if it lands.*) A measurement used in diagnosis or treatment planning can be regulated as a medical device. The honest position — this reports the clinician's measurement and marks its own estimates as estimates — is a better place to stand than most, but it is not a legal opinion and one should be obtained early.
3. **Therapists already own a goniometer, and it costs \$12.** The sale is the documentation, never the measurement. Lead with the wrong half and lose.

The one test that settles it

Sit with one solo PT for one afternoon. Time how long ROM documentation takes per patient, times patients per day, times days. **If it is under fifteen minutes a day there is no budget here,** however good the product is.

6 — The floor prep argument

Where the floor is out of tolerance, how much leveller it takes, and the priced change order — before the flooring goes down.

This idea was wrong when it was first written, and the correction is the interesting part.

The correction, up front

The original idea was: *a phone that produces an ASTM E1155 floor flatness F-number, replacing an \$8,995 profiler.*

It cannot, and the standard says why. ASTM E1155 is the accepted method for quantifying finished slab quality [STANDARD](#). Research into automating it reports accuracy requirements on the order of $\pm 0.5\text{ mm over }5\text{ m}$ [REVIEW](#). Phone LiDAR is not within an order of magnitude of that, and **no measured figure for phone LiDAR against E1155 exists in this research.**

So the honest answer is: **a phone cannot certify an F-number, and any product claiming it does is lying or has not checked.**

That killed the idea as conceived. What is below is what survived, and it is a better business than the original.

The job that is actually worth doing

An LVP or LVT installer's tolerance is not E1155. It is published by the flooring manufacturer, in these terms:

- **3/16" over 10 feet** [REVIEW](#)
- **1/8" over 6 feet** — "most LVP manufacturers" [REVIEW](#)
- **1/8" over 4 feet** [REVIEW](#)

And the consequence of getting it wrong is specific and expensive: "lippage, hollow spots, popping seams, or full plank failure within 12–24 months, and **most manufacturer warranties will be voided**" [REVIEW](#). Installing over an out-of-tolerance subfloor "usually becomes a 'pay twice' situation" [REVIEW](#).

So there is a real, recurring, high-stakes argument on every flooring job: is this slab flat enough, whose problem is it, and who pays for the leveller? Today that argument is conducted with a 6-foot straightedge, a flashlight and two people's memories, and the installer usually loses it because they have no evidence.

Where it sits

Product	What it is	Price	Source class
Allen Face D-Meter	Floor profiler (E1155-class)	\$8,995	REVIEW

Product	What it is	Price	Source class
Dipstick (Face Construction)	The original profiler	not published; "less than some other profilers"	VENDOR
Allen Face F-Meter	Rolling profiler	not published	VENDOR
Terrestrial laser scan + Geomagic	Scan-to-heat-map workflow	enterprise	REVIEW
Straightedge and a flashlight	What the installer actually uses	\$40	—
This	Deviation map → straightedge list → leveller quantity → priced change order	flat, ~\$39/mo	—

Nothing sits between \$40 and \$8,995. That gap is the entire business, and the buyer in it is not a certifying engineer — it is a flooring installer who needs to win an argument and price a prep.

The wedge

The phone does not certify. It tells you where to put the straightedge.

This is Trueline's carry list — "the four walls worth a tape, ranked" — transplanted exactly. Scan the floor, get a **relative** deviation map, and the app hands back the ranked list of suspect locations. The installer puts a real 6-foot straightedge on each, types the real gap, and the map re-solves around the measured points.

Then the money: measured deviations → volume of self-levelling underlayment at the required thickness → bags → a **priced floor-prep change order**, with the deviation map and the photographs attached as the evidence.

Every number marked `scanned` or `measured`. **And the report carries a printed refusal: this is not an ASTM E1155 test and must not be used as one.** Saying that plainly is what makes it credible to the people who know the difference — and there is no downside, because an installer never needed an F-number.

The four-pattern check

- **A — trusted data.** A relative map that is honest about being relative, in a dispute currently settled by memory.
- **B — whole job.** Straightedge → guess the bags → argue → maybe a change order. Nothing connects.
- **C — sensor → money.** Deviation → leveller volume → bags → priced change order → signed.
- **D — small operator.** A \$8,995 instrument the flooring installer was never going to buy, for a standard they do not need.

What gets built

Trueline's LiDAR capture, plane-fitting, exact integer geometry, provenance, carry list, price book and change-order pipeline. **The change order already exists in Trueline and is described as the lead sales argument** — here it is the

whole product.

New: relative plane-fit and deviation heat map, simulated straightedge across an arbitrary span, and leveller volume by depth.

The gate and the price

- **Free forever** — scan, deviation map, straightedge list, 2 floors kept.
- **Paid** — leveller quantity, priced change order, report, unlimited floors.

\$39/month, flat. This is a narrow tool and it must be priced like one. It is also the most plausible **first paid product** in the set, because the job is small, the argument is universal, and the buyer feels the pain every week.

Where it fails, ranked by likelihood

1. **Relative accuracy may still not reach 1/8" over 6 feet.** (*Most likely, and it is unmeasured.*) Plane-fitting many points can beat single-point accuracy — but by how much is **not known and not claimed here**. If the phone cannot reliably flag a 3 mm deviation, it cannot even rank where to put the straightedge, and there is no product. **Test this before anything else.**
2. **The installer may not want the evidence.** (*Plausible.*) An installer who documents an out-of-tolerance slab has to raise it, and raising it can lose a job. Some will prefer not to know.
3. **Too narrow to sustain a subscription.** A tool used twice a month gets cancelled in the third month. This may want to be a feature of a flooring estimating product rather than a product.

The one test that settles it

Scan one real slab. Straightedge the same slab properly. **Does the app's ranked list contain the spots the straightedge actually found?** It does not need the right number — it needs the right *locations*. One afternoon, one slab, and the idea lives or dies.

7 — Evidence, not opinion

ADA and accessibility barrier surveys: measured clearances and slopes, a ranked barrier list, a priced remediation plan, and a report built to be read in litigation.

The job today

A building owner gets a demand letter alleging ADA barriers. Or they are buying a property, renewing a lease, or trying to get ahead of both. Someone walks the site with a tape, a smart level, a clipboard and a camera, and measures: door clearances, counter heights, ramp slopes, cross-slopes, reach ranges, turning radii, parking-stall dimensions.

Almost every one of those is a number with a hard pass/fail threshold, and almost every one of them is currently written on a clipboard and retyped twice — once into a report, once into whatever prices the remediation.

The slope numbers matter most and are the worst to capture: a running slope over 1:12 and a cross-slope over 1:48 are failures, and they are read off a two-foot level held by a person crouching on a ramp.

Where it sits

Product	What it is	Price	Source class
BlueDAG	ADA inspection software suite	subscription not published ; enterprise services from \$125/hr	VENDOR
BlueDAG inspection services	Done-for-you site survey	\$2,500–\$4,500 per site	VENDOR
CASp inspection (California)	Certified Access Specialist survey	~\$1,700 average, small business	REVIEW
iWorQ and general inspection tools	Checklists, AHJ-side	various	VENDOR
Tape, smart level, clipboard	What most surveys actually use	~\$60	—
This	Measured survey → barrier list → priced remediation → litigation-grade report	flat, ~\$79/mo	—

The economics are unusually good. A single survey is billed at **\$1,700–\$4,500**. A tool that saves a solo consultant half a day per survey is trivially worth \$79/month — and unlike most of the ideas in this set, **the buyer already charges by the survey**, so the ROI arithmetic is one line long.

And once again the category leader does not publish a subscription price.

The wedge

A barrier survey is evidence, and evidence has provenance or it is worthless.

This is the strongest fit for Trueline's core idea of the eight — stronger, in one specific way, than Trueline itself. A remodeler's wrong dimension costs a cabinet run. **A barrier survey's wrong dimension gets read out in a deposition.**

So:

- **Every dimension is measured or scanned , printed on the report.** A barrier called out on a scanned number that no tape ever confirmed is an expert witness's afternoon.
- **The slope comes off the device's own inertial sensors**, cross-checked against the scan geometry, with the disagreement between the two reported as a number rather than averaged away — the same move Trueline makes with loop closure, which it "reports rather than smooths away."
- **A barrier the surveyor could not reach or could not measure is a marked category**, not a silent omission. Today it is a silent omission.
- **Then the money:** each barrier → remediation item → priced off the price book → a phased plan the owner can actually execute and budget.

That last step is the whole-job move. Today the survey report ends at "this is non-compliant," and the owner has to go find someone to tell them what fixing it costs. **The barrier list and the priced plan are the same document.**

The four-pattern check

- **A — trusted data.** Litigation-grade, pass/fail against published thresholds, currently zero provenance.
- **B — whole job.** Survey → report → remediation pricing → re-inspection are all separate and all retyped.
- **C — sensor → money.** LiDAR + inclinometer → clearance and slope → barrier → priced remediation → signed report.
- **D — small operator.** A solo accessibility consultant competing against a firm charging \$2,500–\$4,500 a site.

What gets built

Trueline's geometry, provenance, photos-with-pose, offline capture, price book, PDF pipeline and report generator carry over nearly intact. The device's motion sensors give slope. **docs/v3.md already plans re-scan and diff — here that becomes re-inspection after remediation, which is a second billable event on every job and the retention mechanism.**

New: the checkpoint library — 2010 ADA Standards thresholds, and California's CBC Chapter 11B where it differs. **This is a large, versioned, jurisdictional body of rules, and it is the real work.** Smaller than NFPA 25 and better documented, but not small.

The gate and the price

- **Free forever** — measure clearances and slopes against thresholds, 1 site.

- **Paid** — barrier list, priced remediation plan, report, re-inspection diff, unlimited sites.

\$79/month flat. One survey a month at \$1,700 makes this a rounding error to the buyer.

Where it fails, ranked by likelihood

1. **The market may be too small.** (*Most likely.*) CASp is a **California** credential, not a national licence [REVIEW](#).
Outside California the buyer is a loose population of accessibility consultants, architects and facilities people with no register to sell into. **Distribution, not product, is the binding constraint** — and this is the idea in the set where I am least able to size the buyer population from public sources.
2. **Liability by association.** (*Plausible.*) If a survey done in your app misses a barrier and the owner is sued, you will be named. The provenance model is a genuine defence — it records exactly what was and was not verified — but it must be designed as a defence from the first commit.
3. **Standards maintenance never stops**, and a stale threshold is worse than no product.

The one test that settles it

Take one real CASp or ADA barrier report and reproduce it from the app's data model, then ask the consultant who wrote it a single question: "**would you put your name on this version?**" Their hesitation, if any, is the product spec.

8 — What is behind this wall

The pre-drywall record: every pipe, wire, stud and duct captured with real coordinates before the walls close — sold as a deliverable, not a subscription feature.

The one idea in this set that creates a **new billable line item** rather than making an existing one cheaper.

The job today

For about three days on every renovation and every new build, the walls are open and everything is visible. Then they close, and for the next forty years everybody who touches that building is guessing.

The electrician cutting into a wall guesses. The plumber guesses. The next remodeler cuts an exploratory hole. The homeowner hangs a shelf and hits a PEX line.

Today the record of that three-day window is a handful of phone photos in somebody's camera roll, with no scale, no position and no way to find them again in four years. They are effectively lost the moment the job ends.

Where it sits

Product	What it is	Price	Source class
ScanManifold	Photo documentation; Photo+Scan adds 3D measure	\$16/user/mo (Photo), \$24/user/mo (Photo+Scan), no seat minimums	REVIEW
DocuSketch	360° capture + sketch + scope, restoration-focused; hardware included	not published	VENDOR
Raken, Fieldwire, Procore	General construction documentation	per-user, tiered	VENDOR
Camera roll	What actually happens	free	—
This	Positioned pre-drywall record → permanent, transferable building document	per-record + subscription	—

Everybody in this table sells the contractor a subscription to store photos. Nobody sells the *building* a permanent record, and nobody has noticed that the person who benefits most is not the contractor at all.

The wedge

The record is anchored in measured space, and it is sold to the homeowner.

1. **Position, not photographs.** A photo of a stud bay tells you nothing in four years. A photo with the camera's exact pose attached, registered into a measured room model, tells you *the cold water line runs 34 3/8" off the east wall, between 41"*

and 44" up. **Trueline already photographs the room every two seconds with the camera's exact pose attached to each shot.** That capability exists and is currently used only as a visual record — here it is the entire product.

2. What was seen and what is inferred. A pipe visible in three frames and triangulated is **measured** . A pipe that disappears behind a joist and was reasoned about is **derived** , and prints that way. **Somebody is going to cut into this wall based on your document.** The distinction between "we saw it" and "we think it continues" is the difference between a product and a liability.

3. The business model nobody in the table is using. This is not sold as a contractor's storage subscription. It is sold as **a deliverable the contractor charges the homeowner for** — \$150–\$400 a job **ASSUMPTION** , on the same invoice as everything else, at essentially zero cost to produce because the walls were open anyway and the scan takes twenty minutes.

The contractor's pitch to the homeowner writes itself: *"When you sell this house, or the next guy opens this wall, you will have this."* It is a differentiator on the bid, a margin line, and a reason to be chosen over the contractor who did not offer it.

And it produces something genuinely new: **a building record that transfers with the property.** Self-contained, opened with no app and no login — Trueline already ships exactly that artefact as "one self-contained file a homeowner opens with no app and no login."

The four-pattern check

- **A — trusted data.** Seen vs. inferred, on a document somebody will cut a wall open trusting.
- **B — whole job.** Photos in one place, plans in another, as-builts nowhere.
- **C — sensor → money.** LiDAR + posed photos → positioned record → a line item on the invoice. **The only one of the eight where the sensor output is the sellable good.**
- **D — small operator.** A per-user photo subscription is pure cost to a two-man shop. This is revenue.

What gets built

The highest reuse of any idea here except templating:

Needed	State
LiDAR room capture	Built
Photos every 2s with exact camera pose	Built
Exact integer geometry and provenance	Built
Self-contained client file, no app or login	Built
Offline-first	Built
Invoicing and the price book	Built

Needed	State
Register a photo's pixel to a wall coordinate	New — the core work
Trace a run across frames; label services	New
Long-term archive format and property transfer	New, and mostly a format decision

The archive is the one place the no-server rule strains. A record meant to outlive the job cannot depend on a contractor's iCloud. The honest answer is that the deliverable is a **self-contained file the homeowner owns and stores** — which is exactly Trueline's existing pattern, and which keeps marginal cost at zero. A hosted vault is an obvious later tier and should be priced as its own, exactly as `docs/v3.md` argues for the hosted client link.

The gate and the price

- **Free forever** — capture, view, 1 record kept.
- **Paid** — issue the record as a deliverable, unlimited records, branding.

\$49/month, plus the contractor charging \$150–\$400 per record **ASSUMPTION**. **Lead with the second number.** This is the only idea in the set where the pitch is "this makes you money" rather than "this saves you money," and that is a much shorter sale.

Where it fails, ranked by likelihood

1. **The value is real and arrives four years late.** (*Most likely, and it is the structural problem.*) The homeowner pays today for a benefit they may never consciously use. That is the economics of insurance, and insurance is sold, not bought. **If contractors will not pitch it, it does not sell** — and asking a contractor to sell something extra at closing is asking a lot.
2. **Registering a pipe in a photo to a wall coordinate is real computer vision.** (*Plausible.*) The pose is known, which helps enormously. Whether a person can tap a pipe in one frame and get a trustworthy 3D position across frames is **unmeasured**. The safe degradation — tap the pipe in two frames from different angles and triangulate manually — is slower and probably sufficient. Build that first.
3. **Liability, and it is worse than it looks.** Somebody will cut into a wall trusting this document and hit something. The `measured / derived` marking is the defence, and the record must say plainly that it is a record, not a locator, and never replaces scanning the wall before cutting.

The one test that settles it

Ask five remodelers: "if this took twenty minutes with the walls open, would you put \$250 on the invoice for it?" **You are testing whether they will sell it, not whether they would use it.** If four say yes, this is the best business in the set.

Sources

Every URL consulted, 2026-08-27, grouped by the brief that uses it.

Provenance classes, as used throughout:

- **VENDOR** — read off the company's own page.
- **REVIEW** — reported by a third-party review, comparison or industry site, **not independently confirmed**.
- **STANDARD** — a published technical or legal standard.
- **ASSUMPTION** — mine, marked at the point of use.

A warning about the **REVIEW** class

A large share of the pricing below comes from comparison and "best X software 2026" sites. **Many of these are operated by a competitor in the same category, or are affiliate-monetised.** Two examples visible in this research: the Yembo-vs-HomeSurvey.ai comparison is published by HomeSurvey.ai, and a PT EMR pricing comparison naming WebPT, Prompt and Raintree is published by SPRY, a competitor to all three.

Treat every **REVIEW figure as a lead to verify with the vendor, not a fact.** Before acting on any of these numbers, get them from the vendor directly.

1 — Templating

- <https://slabwise.com/guide/templating-comparison>
- <https://slabwise.com/best/laser-templating>
- <https://slabwise.com/reviews/proliner-review>
- <https://www.prodim-systems.com/industries/stone-industry-solutions/>
- <https://www.prodim-systems.com/digital-templating-technology-proliner-vs-laser/>
- <https://stonemachinerylocator.com/> (used LT-55 listings)
- <https://alternativeto.net/software/polycam--lidar-3d-scanner/>

2 — Fire & life-safety inspection

- <https://www.inspectpoint.com/>
- <https://www.inspectpoint.com/trades/fire-sprinkler/>
- <https://www.inspectpoint.com/best-fire-inspection-software/>
- <https://www.buildingreports.com/>
- <https://iworq.com/systems/fire-inspections/>
- <https://safetyculture.com/apps/fire-sprinkler-inspection-software>

- <https://fireinspecthub.com/guides/best-fire-inspection-software/>

3 — The binding estimate (moving)

- <https://yembo.ai/moving>
- <https://yembo.ai/moving/ai-surveys>
- <https://yembo.ai/moving/estimates>
- <https://homesurvey.ai/compare/homesurvey-vs-yembo/> — *published by a Yembo competitor*
- <https://www.smartmoving.com/pricing>
- <https://www.moveitpro.com/pricing>
- <https://www.capterra.com/compare/154185-180048/MoveitPro-vs-SmartMoving-Software>
- <https://moveblueprint.com/virtual-survey-software>
- <https://www.atlasvanlines.ca/moving-cost-estimator/>

4 — The honest weight (livestock)

- <https://cattleweightestimation.com/>
- <https://luzardo.tech/>
- <https://mjelivestockequipment.com/cattlenworth/>
- <https://apps.apple.com/us/app/cattle-weight-pro/id6747881785>
- <https://www.marubeni.com/en/news/2025/info/00035.html>
- <https://tuul.ai/research/cattle-weight-estimation>
- <https://www.centralcarolinascale.com/cattle-scales.htm>
- <https://agriid.com/products/livestock-scales>
- <https://www.cattlescales.com/collections/tru-test-scales>

5 — Range of motion

- <https://www.sprypt.com/blog/2026-pt-emr-pricing-spry-vs-webpt-vs-prompt-vs-raintree> — *published by a competitor to the products it prices*
- <https://ehrinsider.com/reviews/jane-app/>
- <https://www.choosingtherapy.com/jane-app-review/>
- <https://patientstudio.com/blog/jane-app-alternatives-for-physical-therapy-practices>
- <https://halomedicaldevices.com/blog/measuring-range-of-motion-physical-therapists/>
- <https://crimsonpublishers.com/ggs/fulltext/GGS.000610.php> — goniometry app accuracy
- <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC6787724/> — smartphone craniocervical ROM validity
- <https://hemscap.com/blogs/167/goniometry-measuring-joint-range-of-motion-in-physical-therapy>

6 — Floor prep

- <https://allenface.com/d-meter-floor-profiler/> — D-Meter at \$8,995
- <https://allenface.com/d-meter-floor-profiler/d-meter-floor-profiler-vs-dipstick/>
- <https://allenface.com/f-meter-rolling-floor-profiler/>
- <https://dipstick.com/>
- <https://ifti.com/astm-e1155-testing-standard/> **STANDARD**
- <https://ifti.com/floor-flatness-and-levelness-testing-complete-astm-e1155-procedure-guide/> **STANDARD**
- <https://www.archtoolbox.com/ff-and-fl-numbers/>
- <https://www.concretenetwork.com/industrial-floors/flatness.html>
- <https://www.researchgate.net/publication/268466745> — E1155 sampling requirements
- <https://www.perfectfloorsinc.com/blog/articles/subfloor-flatness-101-the-tolerance-that-determines-if-your-lvp-click-system-fails>
- <https://elliecm.com/lvt-flooring-how-flat-is-flat-enough/>
- <https://brokeringsolutions.com/wp-content/uploads/2021/12/LVT-Flooring-Standard-Installation-Guide-08-05-21.pdf> **VENDOR**
- <https://truinstallation.com/blog/how-to-prepare-subfloor-for-lvp-installation>

7 — Barrier survey

- <https://www.bluedag.com/packages.html>
- <https://www.bluedag.com/inspections.html>
- <https://www.bluedag.com/software/>
- <https://www.adacertified.com/casp-inspection-cost>
- <https://adacasp.com/casp-inspection-california/>
- <https://caspinspectors.com/>
- <https://calaccessibility.com/>

8 — What's behind this wall

- <https://www.scanmanifold.com/blog-posts/pre-drywall-inspection-checklist-contractors-2026>
 - <https://www.scanmanifold.com/blog-posts/pre-close-inspection-checklist-before-drywall-2026>
 - <https://www.scanmanifold.com/blog-posts/renovation-contractor-photo-documentation-2026>
 - <https://www.g2.com/products/docusketch/reviews>
 - <https://contractortoolstack.com/categories/photo-documentation/>
-

Figures deliberately NOT claimed

Recorded here so nobody later mistakes a gap for an oversight:

Figure	Status
Phone LiDAR accuracy for countertop templating (± 1.5 mm)	Not found, not measured. Decides idea 1.
Phone LiDAR relative accuracy at 1/8" over 6 ft	Not found, not measured. Decides idea 6.
Phone LiDAR against ASTM E1155 (± 0.5 mm / 5 m)	Established as out of reach. Idea 6 was rewritten because of it.
On-device pose estimation vs. goniometry	Not found, not measured. Decides idea 5.
Independent validation of any cattle weight app's accuracy claim	Not found. Is the wedge for idea 4.
Inspect Point, BuildingReports, BlueDAG, DocuSketch, Yembo subscription pricing	Not published by the vendor. Quote-only.
Number of accessibility consultants outside California	Not found. Is the main risk to idea 7.
What a homeowner will pay for a pre-drywall record	ASSUMPTION — \$150–\$400. Untested.